



Εθνικό Μετσόβιο Πολυτεχνείο
ΔΠΜΣ Επιστήμη Δεδομένων και Μηχανική Μάθηση

**ONASSIS
FOUNDATION**

Accelerating SIVIA (Set Inversion via Interval Analysis). An Interval Set Membership Technique to explain Neural Classifier Decisions.

Nasiotis Konstantinos

Contents

Contents

- ❖ Introduction
 - ❖ Background
 - ❖ Implementation
 - ❖ Results
 - ❖ Conclusion
 - ❖ Acknowledgements
-

Introduction

Interval Analysis

What is Interval Analysis?

- Interval Analysis is a field of mathematics which was created by R. E. Moore in 1962 with the purpose of integrating errors in engineering computational problems

What are some common applications of IA?

- Basically, any engineering problem can benefit from Interval Analysis.
- Applications range from path planning and global optimization to neural networks

What is SIVIA?

Set Inversion Via Interval Analysis is a technique utilizing interval analysis that can estimate the input space of non-linear functions

SIVIA Applications

❖ Parameter Estimation

- ❖ I. Braems, F. Berthier, L. Jaulin, M. Kieffer, and E. Walter, *Guaranteed estimation of electrochemical parameters by set inversion using interval analysis*, *Journal of Electroanalytical Chemistry*, vol. 495, no. 1, pp. 1–9, 2000.

❖ Combinatorial & Global Optimization problems

- ❖ E. Hansen and G. W. Walster, *Global optimization using interval analysis: revised and expanded*, vol. 264. CRC Press, 2003.

❖ Path Planning

- ❖ A. Pruski and S. Rohmer, *Robust path planning for non-holonomic robots*, *Journal of Intelligent and Robotic Systems*, vol. 18, pp. 329–350, 1997.
- ❖ L. Jaulin, “Path planning using intervals and graphs,” *Reliable computing*, vol. 7, no. 1, pp. 1–15, 2001.

❖ Neural Networks

- ❖ Vladik Kreinovich and Andrew Bernat. *Parallel algorithms for interval computations: An introduction*. Interval Computations, 1994, 01 1994.
 - ❖ S. P. Adam, D. A. Karras, G. D. Magoulas, and M. N. Vrahatis, *Reliable estimation of a neural network’s domain of validity through interval analysis-based inversion*, in *2015 International Joint Conference on Neural Networks (IJCNN)*, pp. 1–8, 2015.
 - ❖ S. P. Adam, A. C. Likas, and M. N. Vrahatis, *Evaluating generalization through interval-based neural network inversion*, *Neural Computing and Applications*, vol. 31, no. 12, pp. 9241–9260, 2019.
-

Interval Analysis

What are the challenges of SIVIA?

- SIVIA works by exhaustively expanding the search space, increasing computational demand exponentially

Why acceleration?

- Parallelization can reduce computational intensity by distributing work to other processing units, in this case, GPU devices

Background

Interval Definitions

Interval $[x] \subset \mathbb{R}$

- Lower Bound $lb([x]) = \underline{x}$
- Upper Bound $ub([x]) = \bar{x}$
- Diameter $Width([x]) = |\bar{x} - \underline{x}|$
- Midpoint (or Center) $Mid([x]) = \frac{\underline{x} + \bar{x}}{2}$

Example:

- $[x] = [\underline{x}, \bar{x}] = [3, 5]$
 - $Width([x]) = 5 - 3 = 2$
 - $Mid([x]) = \frac{3+5}{2} = 4$
-

Interval Basic Operations

- Addition

$$[x] + [y] = [\underline{x} + \underline{y}, \bar{x} + \bar{y}]$$

e.g. $[3,5] + [2,5] = [5,10]$

- Subtraction

$$[x] - [y] = [\underline{x} - \bar{y}, \bar{x} - \underline{y}]$$

e.g. $[3,5] - [2,5] = [-2,3]$

- Multiplication

$$[x] * [y] = [\min\{\underline{x}\underline{y}, \underline{x}\bar{y}, \bar{x}\underline{y}, \bar{x}\bar{y}\}, \max\{\underline{x}\underline{y}, \underline{x}\bar{y}, \bar{x}\underline{y}, \bar{x}\bar{y}\}]$$

e.g. $[3,5] * [2,5] = [6,25]$

- Division

$$\frac{[x]}{[y]} = [\underline{x}, \bar{x}] * \frac{1}{[\underline{y}, \bar{y}]} \text{ where } \frac{1}{[\underline{y}, \bar{y}]} = \left[\frac{1}{\bar{y}}, \frac{1}{\underline{y}}\right] \text{ if } 0 \notin [\underline{y}, \bar{y}]$$

e.g. $\frac{[3,5]}{[2,5]} = [0.6, 2.5]$

Interval Set Operations

Set-Membership Operations

- isSubset

$$[x] \subseteq [y] = \underline{y} \leq \underline{x} \text{ AND } \bar{y} \geq \bar{x}$$

- Intersects

$$[x] \cap [y] = (\underline{x} \leq \bar{y} \text{ AND } \bar{x} \geq \underline{y})$$

Interval Trigonometric Operations


Non-linear Elementary & Trigonometric functions are also supported!

- $[x]^2$
 - $\sqrt{[x]}$
 - $e^{[x]}$
 - $\log([x])$
 - $\sin([x])$
 - $\cos([x])$
 - $\tan([x])$
-

Interval Boxes

- Multidimensional Intervals or Interval Vectors

$$[X] = [[x_1], [x_2], \dots, [x_n]]$$




Interval Matrices also supported!!

Interval Functions

- Interval Inclusion Functions

Given f from \mathbb{R}^n to \mathbb{R}^m , $[f]$ (from \mathbb{IR}^n to \mathbb{IR}^m) is an inclusion function if

$$\forall [x] \in \mathbb{IR}^n \quad f([x]) \subset [f]([x])$$



SIVIA Algorithm

SIVIA Algorithm

Input: X_0 Box, incl. function $[f](x)$, image Y

- Begin with an initial X_0 box.

Repeat:

- Send Box to an Inclusion Function $[f]([x])$
- Evaluation phase - $[f]([x]) \subseteq Y$
- **Accept** - $[f]([x]).\text{isSubset}(Y)$ (*Box is part of the solution*)
- **Reject** - $\neg [f]([x]).\text{intersects}(Y)$
- **Bisect** - $X_0 \rightarrow X_1, X_2$
- Also reject $\text{width}(\text{box}) \leq \epsilon$ (*Box is too small to process*)

Until all boxes have been evaluated.


Output: X_n Boxes with their categorization labels (accepted, rejected, epsilon)

Branch and Bound Technique!!

❖ Jaulin L., Kieffer M., Didrit O., Walter E. (2001). *Applied interval analysis*. Springer.

❖ L. Jaulin and E. Walter, *Set inversion via interval analysis for nonlinear bounded-error estimation*, *Automatica*, vol. 29, no. 4, pp. 1053–1064, 1993.

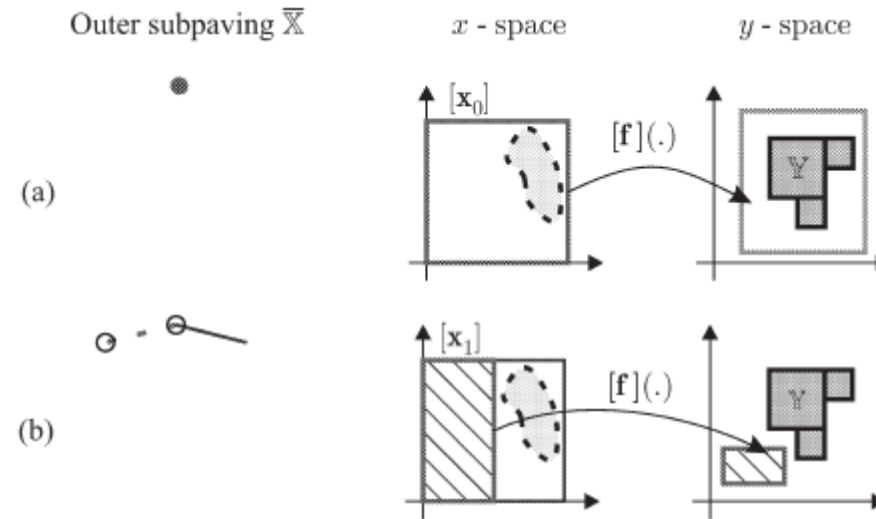
SIVIA Algorithm



❖ Jaulin L., Kieffer M., Didrit O., Walter E. (2001). *Applied interval analysis*. Springer.

❖ L. Jaulin and E. Walter, *Set inversion via interval analysis for nonlinear bounded-error estimation*, *Automatica*, vol. 29, no. 4, pp. 1053–1064, 1993.

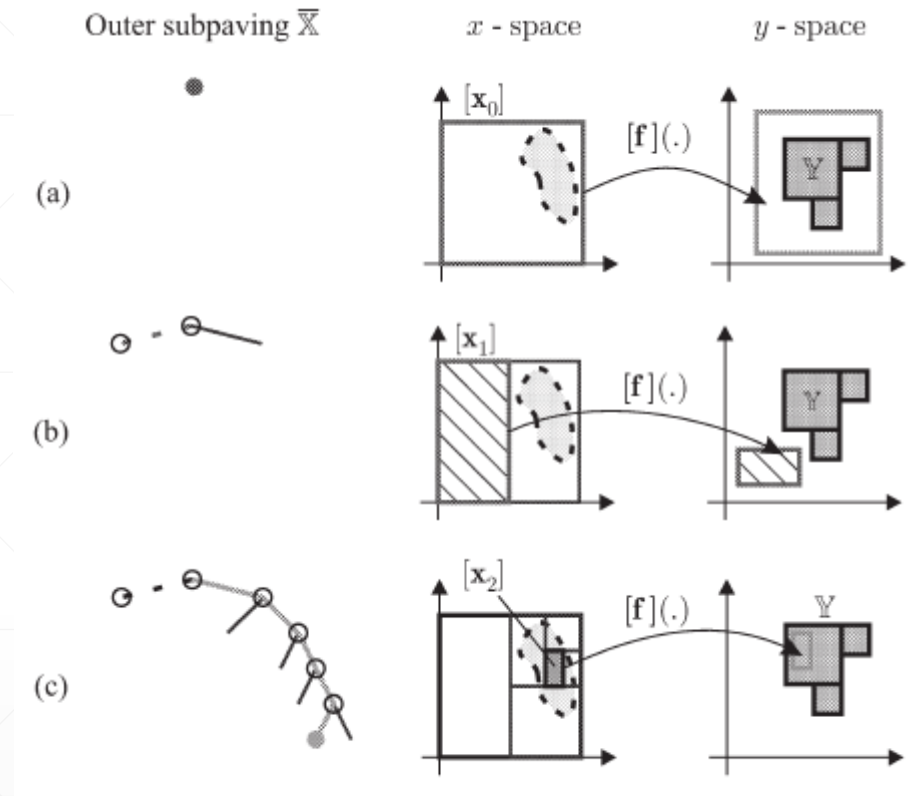
SIVIA Algorithm



❖ Jaulin L., Kieffer M., Didrit O., Walter E. (2001). *Applied interval analysis*. Springer.

❖ L. Jaulin and E. Walter, *Set inversion via interval analysis for nonlinear bounded-error estimation*, *Automatica*, vol. 29, no. 4, pp. 1053–1064, 1993.

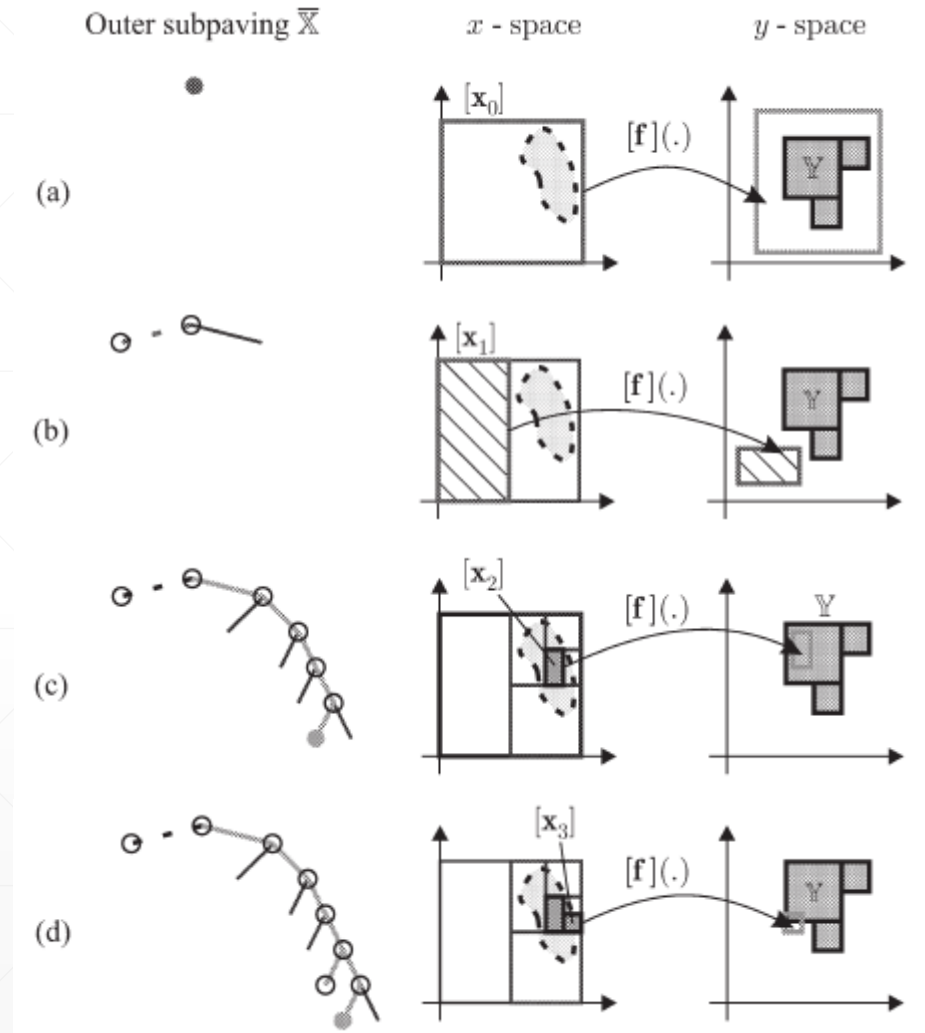
SIVIA Algorithm



❖ Jaulin L., Kieffer M., Didrit O., Walter E. (2001). *Applied interval analysis*. Springer.

❖ L. Jaulin and E. Walter, *Set inversion via interval analysis for nonlinear bounded-error estimation*, *Automatica*, vol. 29, no. 4, pp. 1053–1064, 1993.

SIVIA Algorithm






❖ Jaulin L., Kieffer M., Didrit O., Walter E. (2001). *Applied interval analysis*. Springer.

❖ L. Jaulin and E. Walter, *Set inversion via interval analysis for nonlinear bounded-error estimation*, *Automatica*, vol. 29, no. 4, pp. 1053–1064, 1993.

SIVIA and Neural Networks

SIVIA & Neural Networks

- ❖ SIVIA enables the estimation of non-linear spaces
- ❖ Consequently, combining it with a Neural Network enables the estimation of its input space even without knowledge of training data
- ❖ Allows the extraction of useful information such as the volume of a recognized input area



- ❖ McCulloch, Warren S., and Walter Pitts. *A logical calculus of the ideas immanent in nervous activity*. *The bulletin of mathematical biophysics* 5 (1943): 115-133.
- ❖ S. Haykin, *Neural Networks and Learning Machines*. Pearson India, 2008.
- ❖ S. P. Adam, A. C. Likas, and M. N. Vrahatis, "Evaluating generalization through interval-based neural network inversion," *Neural Computing and Applications*, vol. 31, no. 12, pp. 9241–9260, 2019.

GPU Parallelism

GPU Parallelism


- ❖ SIMD/SIMT: many low performance cores solving the same task in parallel
- ❖ Originally developed for computer graphics
- ❖ Primary focus on NVIDIA GPUs due to high market share



- ❖ Lorenz Gillner, Ekaterina Auer., *Interval Methods for the GPU, SWIM*, 2023.
- ❖ P. Pacheco and M. Malensek, *An introduction to parallel programming*. Morgan Kaufmann, 2021
- ❖ <https://docs.nvidia.com/cuda/cuda-c-programming-guide/>

CUDA Programming Model

- ❖ Defined by Blocks, Threads, Warps
- ❖ Blocks are programmable
- ❖ Warps are groups of 32 Threads
- ❖ A thread is a processing (CUDA) core (max 1024 per block)



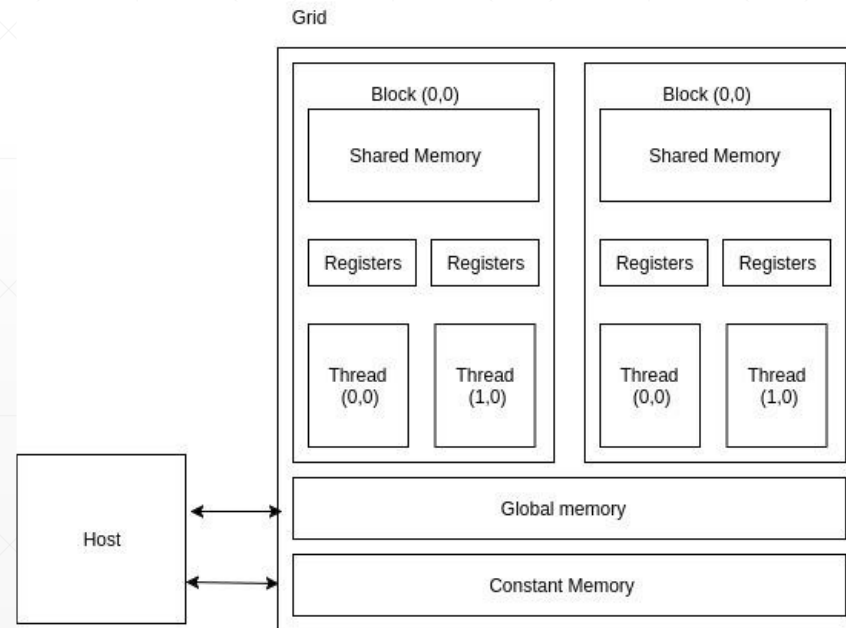
❖ Lorenz Gillner, Ekaterina Auer., *Interval Methods for the GPU, SWIM*, 2023.

❖ P. Pacheco and M. Malensek, *An introduction to parallel programming*. Morgan Kaufmann, 2021

❖ <https://docs.nvidia.com/cuda/cuda-c-programming-guide/>

CUDA Memory Hierarchy

- ❖ Constant, Global, Shared, Registers
- ❖ Global Memory is the slowest but largest in size – also known as VRAM
- ❖ Registers are the fastest but very small in size
- ❖ Shared Memory is on-chip slower than registers but of larger size
- ❖ Cache reduces the cost of accessing the Global Memory



Implementation

CUDA Design Principles

- ❖ Find ways to parallelize sequential code.
- ❖ Minimize data transfers between the host and the device.
- ❖ Adjust kernel launch configuration to maximize device utilization.
- ❖ Ensure global memory accesses are coalesced.
- ❖ Minimize redundant accesses to global memory whenever possible.
- ❖ Avoid long sequences of diverged execution by threads within the same warp.

Implementation

Single GPU:


- X_0 Box initialization.
- Initialize work pool on HOST & DEVICE.
- If $\text{calculateBoxes}(X_0) > \text{Capacity}_{\text{device}}$
 - Bisect X_0 on the HOST
- Move X_0 Box to the HOST pool.
- Copy pool from $\text{HOST} \rightarrow \text{DEVICE}$.
- Parallel Bisection.
- Inclusion Function.
- Set-Membership Operations.
- Transfer the result $\text{DEVICE} \rightarrow \text{HOST}$

GPU

GPU

GPU

CPU



Implementation

Single GPU:


- X_0 Box initialization.
- Initialize work pool on HOST & DEVICE.
- If $\text{calculateBoxes}(X_0) > \text{Capacity}_{\text{device}}$
 - Bisect X_0 on the HOST
- Move X_0 Box to the HOST pool.
- Copy pool from $\text{HOST} \rightarrow \text{DEVICE}$.
- Parallel Bisection.
- Inclusion Function.
- Set-Membership Operations.
- Transfer the result $\text{DEVICE} \rightarrow \text{HOST}$

GPU

GPU

GPU

CPU



SSP Strategy!!

- ❖ B. Gendron and T. G. Crainic, *Parallel branch-and-branch algorithms: Survey and synthesis*, Operations research, vol. 42, no. 6, pp. 1042–1066, 1994.
- ❖ K. Nasiotis, D. López, S. Adam, and L. Casado, *Set inversion via interval analysis a study on parallel processing implementation*, SWIM 2019.

Implementation

Multiple GPUs:

- X_0 Box initialization.
- Initialize work pools on HOST & DEVICES.
- Bisect X_0 on the HOST
- Copy pools from $HOST \rightarrow DEVICES$.
- Start threads equal to number of DEVICES
- If $\text{calculateBoxes}(X_{thread}) > \text{Capacity}_{device}$
 - Bisect X_{thread}
- Copy Buffers to each DEVICE Buffer
- Parallel Bisection.
- Inclusion Function.
- Set-Membership Operations
- Transfer the result.

Thread

Thread


Thread

GPU

GPU

GPU

Thread



AMP/SSP Hybrid Strategy!!

- ❖ B. Gendron and T. G. Crainic, *Parallel branch-and-branch algorithms: Survey and synthesis*, Operations research, vol. 42, no. 6, pp. 1042–1066, 1994.
- ❖ K. Nasiotis, D. López, S. Adam, and L. Casado, *Set inversion via interval analysis a study on parallel processing implementation*, SWIM 2019.

Parallel Bisection

Input:


- Pool with X_0 Box at the 0th index

Algorithm:

- Host sends a number **bisect** commands to the DEVICE.
- Number of commands determined by $\log_2(\text{numBoxes})$

Output:

- Complete pool of N Boxes.



Parallel Bisection

Pros:

- ❖ Faster Execution
- ❖ Trivial memory transfer costs
- ❖ Scalable

Cons:


- ❖ Wasted GPU Resources in initial executions

```
> GPU Device has Compute Capabilities SM 8.6


Bisection Benchmark
Initial Box ([-1.500000,1.500000],[-1.500000,1.500000])
Epsilon: 0.001 Dimensions: 2
Problem size: 16777216 boxes
Boxes: 1 | CPU Duration: 0us GPU Duration: 230us
Boxes: 2 | CPU Duration: 0us GPU Duration: 40us
Boxes: 4 | CPU Duration: 0us GPU Duration: 43us
Boxes: 8 | CPU Duration: 0us GPU Duration: 49us
Boxes: 16 | CPU Duration: 0us GPU Duration: 44us
Boxes: 32 | CPU Duration: 1us GPU Duration: 43us
Boxes: 64 | CPU Duration: 2us GPU Duration: 54us
Boxes: 128 | CPU Duration: 5us GPU Duration: 47us
Boxes: 256 | CPU Duration: 9us GPU Duration: 75us
Boxes: 512 | CPU Duration: 19us GPU Duration: 48us
Boxes: 1024 | CPU Duration: 39us GPU Duration: 43us
Boxes: 2048 | CPU Duration: 77us GPU Duration: 40us
Boxes: 4096 | CPU Duration: 216us GPU Duration: 40us
Boxes: 8192 | CPU Duration: 451us GPU Duration: 127us
Boxes: 16384 | CPU Duration: 916us GPU Duration: 43us
Boxes: 32768 | CPU Duration: 1693us GPU Duration: 43us
Boxes: 65536 | CPU Duration: 3398us GPU Duration: 46us
Boxes: 131072 | CPU Duration: 6843us GPU Duration: 56us
Boxes: 262144 | CPU Duration: 13761us GPU Duration: 201us
Boxes: 524288 | CPU Duration: 28239us GPU Duration: 236us
Boxes: 1048576 | CPU Duration: 55536us GPU Duration: 307us
Boxes: 2097152 | CPU Duration: 111522us GPU Duration: 433us
Boxes: 4194304 | CPU Duration: 225368us GPU Duration: 688us
```

Parallel Bisection


Equal Bisection \rightarrow Full expansion of the problem tree




$e = 0.1$



$e = 0.01$



$e = 0.4$



$e = 0.2$

Inference & Set Estimation

Input:

- Vector of Boxes – Work Pool
- Empty Label Vector - Problems 1&2

Algorithm:

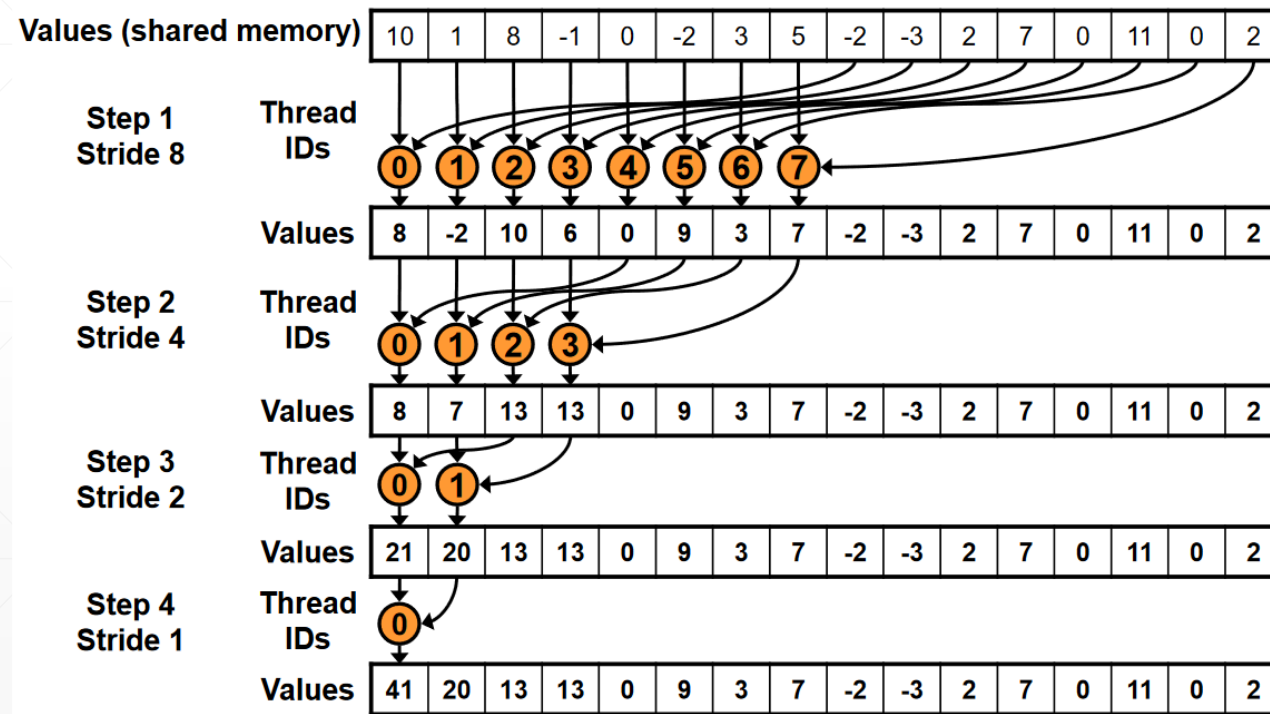
- Coarse-grained approach
- Each Box processed independently. ($[f]([x])$ and set-membership operations)
- Optimized to avoid Warp Divergence. (replaces “if” clauses with Boolean operations)
- Parallel Reduction – Problem 3

Output:

- Label Buffer filled with evaluation labels (included, discarded, epsilon) – Problems 1&2
 - Volume of domain of validity – Problem 3
-

Parallel Reduction

- ❖ Without Warp Divergence
- ❖ Sequential Addressing (cache friendly)





Results

Problems 1 & 2



❖ Problem 1: 2D Torus

$$\begin{aligned} [f]([x]) &= [x]^2 + [y]^2 \\ [x]_0 &= [-1.5, 1.5], [-1.5, 1.5] \\ [y] &= [1, 2] \end{aligned}$$



❖ Problem 2: 2D Griewank

$$\begin{aligned} [x_0] &= [-10, 10]^2 \\ [y] &= [1.5, 3] \\ [f]([x]) &= \sum_{i=1}^2 \frac{[x]_i^2}{4000} - \prod_{i=1}^2 \cos\left(\frac{[x]_i}{\sqrt{i}}\right) + 1 \end{aligned}$$



- Problem Requirement: *The whole input space*

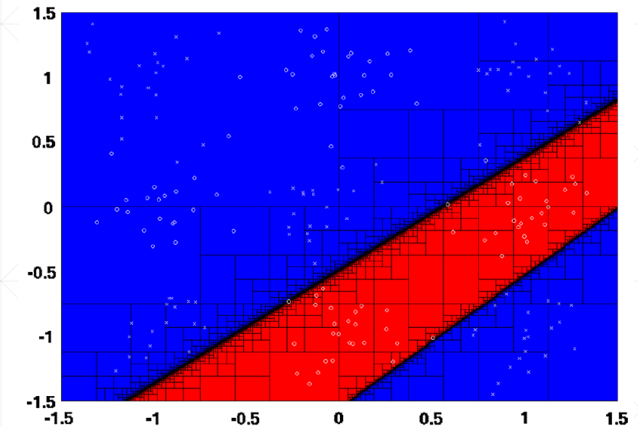
Problem 3

❖ Evaluating Generalization Performance of Neural Classifiers


- $V_{net} = \sum_{i=1}^M V_i$
- $G_{net} = \frac{V_{net}}{V_{input}} - \frac{l}{P}$
- $V_{input} = \prod_i^N |x_i^{max} - x_i^{min}|$
- **Problem Requirement: A single value**

Consequently:


- ❖ $[f]([x]): 6 - 30 - 2 \text{ MLP}$ Weights provided by [1], Vertebral Column Dataset, example with early stopping
- ❖ $[x]_0 = [-1, 1]^6$
- ❖ $[y] = [0.8, 1]$




Undertraining



Sample data



“Proper” fit



Overtraining (overfitting)

- ❖ [1] S. P. Adam, A. C. Likas, and M. N. Vrahatis, “Evaluating generalization through interval-based neural network inversion,” *Neural Computing and Applications*, vol. 31, no. 12, pp. 9241–9260, 2019.
- ❖ S. P. Adam, D. A. Karras, G. D. Magoulas, and M. N. Vrahatis, *Reliable estimation of a neural network’s domain of validity through interval analysis-based inversion*, in 2015 *International Joint Conference on Neural Networks (IJCNN)*, pp. 1–8, 2015.

Test Configurations

Home Setup

- GPU: Nvidia RTX 3060 Ti 8GB
- OS: WSL2 Ubuntu 22.04.2 LTS

Google Collab:

- GPU 1: Nvidia A100 40GB
- GPU 2: Nvidia Tesla V100 16GB

Microlab DaVinci of NTUA:

- GPU: Nvidia Tesla V100 32GB
- OS: Ubuntu 20.04.2 LTS

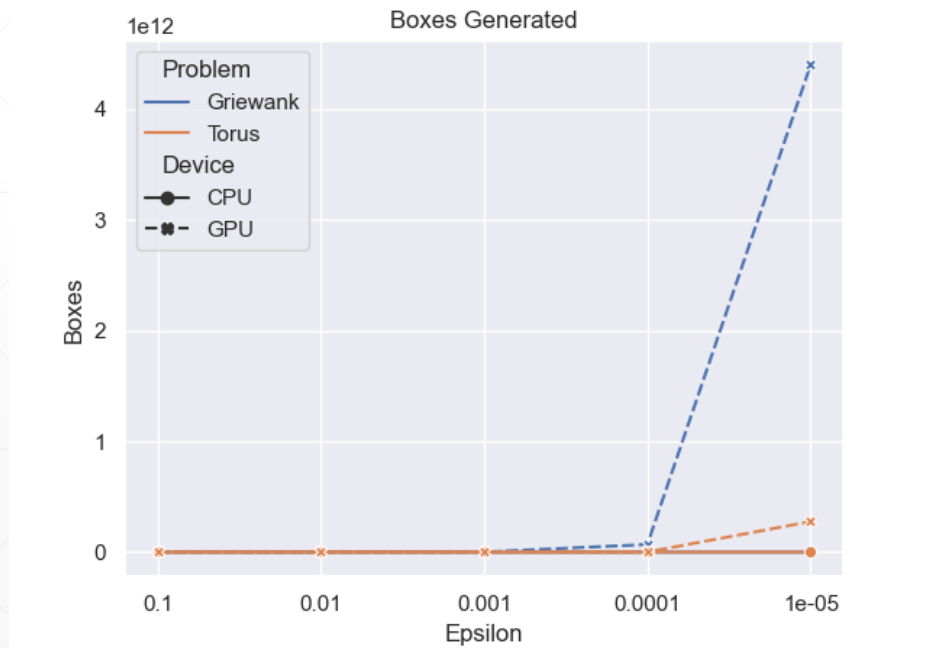
CEID of UPATRAS:

- GPU: 8x Nvidia A100 32GB
 - OS: Ubuntu 22.04.3 LTS
-

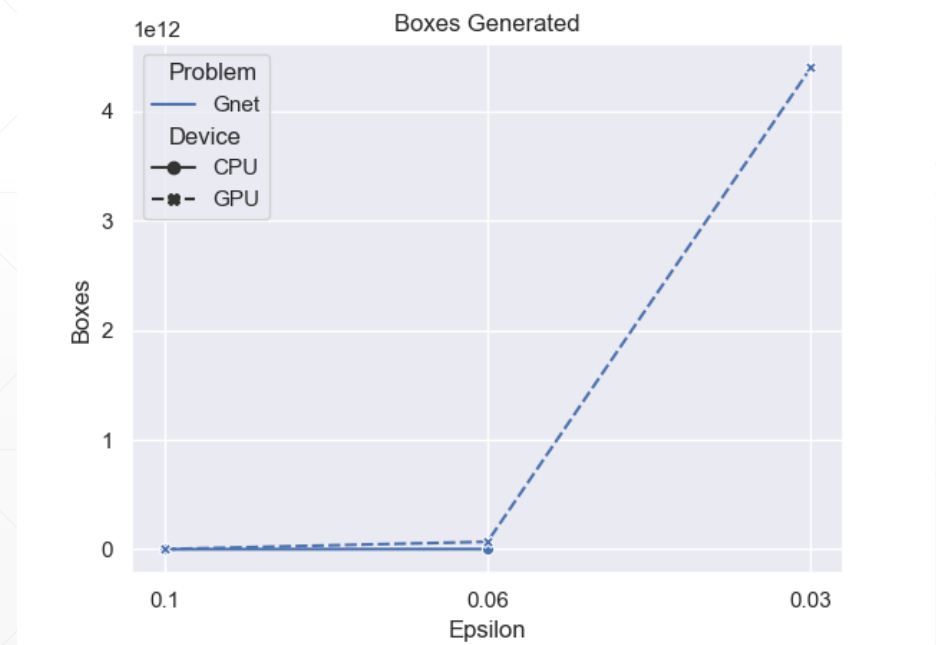
Problem Size

- ❖ Parallel implementation computes a larger number of boxes
- ❖ Missing value on problem 3 due to long execution time required to exit execution

Problems 1&2





Problem 3



Results: Problem 1



Speedup

- ❖ Comparison to sequential SIVIA using an i7 5820k @ 4.3Ghz
 - Slowdowns due to intensive memory transfers between host and device



Throughput


- ❖ CPU has way smaller throughput
- ❖ Half variables increase throughput
 - ❖ Slight depiction due to logarithmical scaling




Kernel Runs

❖ Problem 1

❖ More VRAM -> More Boxes fit -> less GPU executions



❖ FP32





❖ FP16

Results: Problem 2



Speedup

❖ Similar findings as with Problem 1



Throughput


- ❖ Throughput findings of Problem 1 also extend to Problem 2
- ❖ Half variables increase throughput




Kernel Runs

❖ Problem 2

❖ Similar behavior to Problem 1, more VRAM mean less GPU executions



❖ FP32





❖ FP16

Problem 3

Speedup



❖ **Key result:** 5000+ speedup when using 8x A100 GPUs

❖ 8000 with FP16 variables



Throughput


- ❖ Speedup findings align with the throughput
- ❖ CPU throughput shrinks with increasing problem size




Kernel Runs

❖ Problem 3

❖ Like problems 1&2, less VRAM = More executions



❖ FP32



❖ FP16

Profiling

Profiling

❖ Nvidia profiler validates assumptions of previous results

❖ Problem 2
❖ $e = 0.001$

| Time (%) | Total Time (ns) | Num Calls | Avg (ns) | Med (ns) | Min (ns) | Max (ns) | StdDev (ns) | Name |
|----------|-----------------|-----------|-------------|-------------|-----------|------------|--------------|-----------------------|
| 85.2 | 5697886176 | 9 | 633098464.0 | 104344400.0 | 56500 | 3964895784 | 1317343026.4 | cudaMemcpy |
| 10.8 | 722301100 | 1 | 722301100.0 | 722301100.0 | 722301100 | 722301100 | 0.0 | cudaMallocHost |
| 2.5 | 168768499 | 60 | 2812808.3 | 43000.0 | 20200 | 51632499 | 9549282.8 | cudaDeviceSynchronize |
| 0.6 | 43053000 | 3 | 14351000.0 | 4908400.0 | 293300 | 37851300 | 20482257.3 | cudaFree |
| 0.6 | 42401600 | 3 | 14133866.7 | 11234500.0 | 550600 | 30616500 | 15241205.0 | cudaMalloc |

❖ Problem 3
❖ $e = 0.06$

| ** CUDA API Summary (cuda_api_sum): | | | | | | | | |
|-------------------------------------|-----------------|-----------|-------------|-------------|-----------|-----------|-------------|-----------------------|
| Time (%) | Total Time (ns) | Num Calls | Avg (ns) | Med (ns) | Min (ns) | Max (ns) | StdDev (ns) | Name |
| 99.1 | 114912155468 | 7681 | 14960572.3 | 54200.0 | 2600 | 374076197 | 59393960.7 | cudaDeviceSynchronize |
| 0.6 | 684587990 | 1 | 684587990.0 | 684587990.0 | 684587990 | 684587990 | 0.0 | cudaMallocHost |
| 0.2 | 217664907 | 7680 | 28341.8 | 12800.0 | 6300 | 4600899 | 62991.5 | cudaLaunchKernel |
| 0.1 | 58488800 | 10 | 5848880.0 | 3500.0 | 2100 | 58446300 | 18480850.2 | cudaFree |
| 0.0 | 44460599 | 40 | 1111515.0 | 2050.0 | 1600 | 43923599 | 6943062.0 | cudaMalloc |
| 0.0 | 24409100 | 549 | 44461.0 | 40600.0 | 9500 | 317700 | 34166.6 | cudaMemcpy |

Conclusion

Conclusion

- ❖ GPU parallelization is very compatible with SIVIA when intensive memory transfers are not a requirement
 - ❖ Mathematical Techniques should accompany parallelization efforts
 - ❖ E.g. Optimization problems have their search space reduced by calculating the hull of an iteration's *best* Boxes
 - ❖ Proposed implementation not a one-size-fits-all technique
 - ❖ Scales very well with multiple GPUs
-

Future Work

- Expansion of GPU Interval Ecosystem
- Better Fine-tuning and resource allocation
- Test algorithm with optimization problems
- Experiment with in-GPU synchronization strategies
- Test and compare modern parallel interval environments^{[1][2]}

❖ [1]Lorenz Gillner, Ekaterina Auer., *Interval Methods for the GPU, SWIM, 2023.*

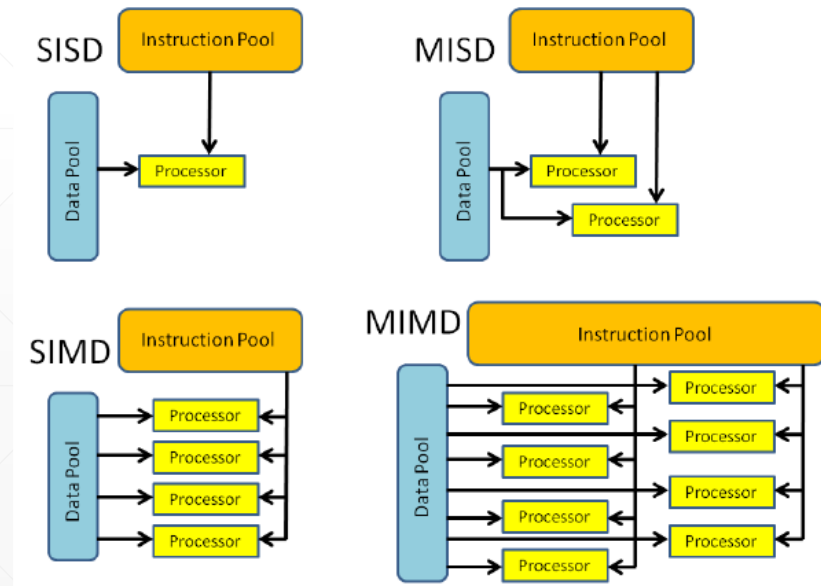
❖ [2]<https://github.com/JuliaIntervals/IntervalArithmetic.jl>

Acknowledgements

Thank you for your time!

Flynn's Taxonomy

- ❖ SISD – Single Instruction Single Data
- ❖ MISD – Multiple Instructions Multiple Data
- ❖ **SIMD** – Single Instruction Multiple Data
- ❖ **MIMD** – Multiple Instructions Multiple Data



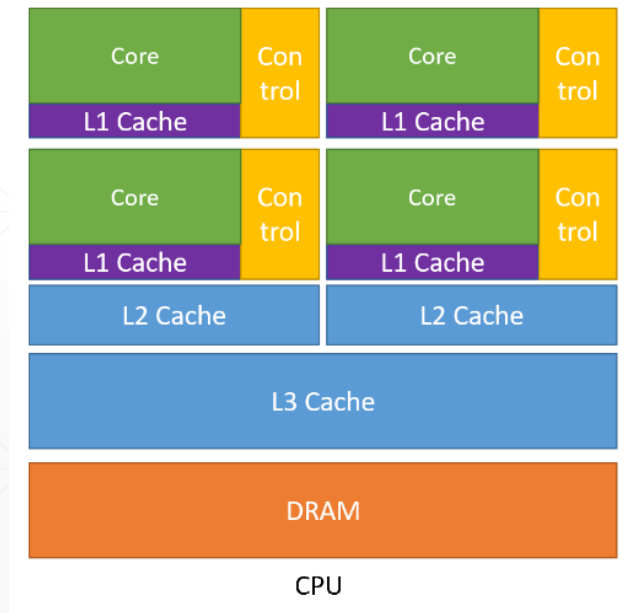
❖ Lorenz Gillner, Ekaterina Auer., *Interval Methods for the GPU, SWIM*, 2023.

❖ P. Pacheco and M. Malensek, *An introduction to parallel programming*. Morgan Kaufmann, 2021

Parallel Computing


❖ Task Parallelism

❖ Data Parallelism



❖ Fine-Grained

❖ Coarse-Grained



❖ Lorenz Gillner, Ekaterina Auer., *Interval Methods for the GPU, SWIM*, 2023.

❖ P. Pacheco and M. Malensek, *An introduction to parallel programming*. Morgan Kaufmann, 2021